

Reconfigurable Energy Efficient Accelerator Architecture for CNN based Edge Applications

Muhammad Usman

*Centres of Excellence in Science &
Applied Technologies (CESAT)
Islamabad, Pakistan
i201320@nu.edu.pk*

Azhar Zahid

*Centres of Excellence in Science &
Applied Technologies (CESAT)
Islamabad, Pakistan
i201312@nu.edu.pk*

Abdul Moiz Sheikh

*Centres of Excellence in Science &
Applied Technologies (CESAT)
Islamabad, Pakistan*

Abstract—Real-time artificial intelligence (AI) applications demand significant computational resources, often exceeding the capabilities of general-purpose processors. Consequently, research focus has shifted towards application-specific hardware, known as AI accelerators. This work presents an optimized AI accelerator architecture emphasizing memory hierarchy design to enhance performance and energy efficiency. Since memory access contributes substantially to overall delay and power consumption, the proposed architecture integrates multi-level memory for efficient data reuse and minimize redundant data transfers. Additionally, a tiling-based processing approach is employed to perform convolution, pooling, activation, and batch normalization operations in parallel across individual tiles, thereby reducing hardware resource utilization. The proposed design achieves a throughput of 90 GOPS with a power consumption of 218 mW after synthesis on SKY-Water 130 nm ASIC technology, corresponding to a power efficiency of 418 GOPS/W at 100 MHz. The results demonstrate a 65% improvement in power efficiency compared to state-of-the-art implementations, validating the effectiveness of the proposed memory-optimized accelerator architecture.

Index Terms—component, formatting, style, styling, insert.

I. INTRODUCTION

Convolutional Neural Networks (CNNs) are among the most widely used architectures in deep learning, particularly suited for processing two-dimensional data through their convolutional layers [1]. From a computational perspective, the convolution operation dominates the workload, involving intensive multiply-accumulate (MAC) operations between input feature maps and kernel weights. These operations require a substantial number of multipliers and adders, as well as frequent memory accesses, leading to significant latency and power consumption [2].

Conventional CPUs lack the parallel computing capability required for such workloads, while GPUs, though offering improved parallelism, suffer from high power consumption and limited power efficiency for real-time AI applications [3]. Consequently, there is a growing demand for dedicated and energy-efficient hardware accelerators capable of supporting diverse AI workloads. Google's Tensor Processing Unit (TPU) represents one such advancement; however, still AI accelerator design remains an active and evolving research area.

Various optimization techniques have been proposed to enhance accelerator performance, including systolic arrays,

memory reuse, loop unrolling, and approximate multipliers. The fundamental computational block in most of these designs is the Processing Element (PE) [4], which performs the MAC operations utilized in CNN processing. Additionally, caching and data reuse are also critical for reducing off-chip memory communication and improving throughput [2].

The motivation behind this research is to develop a custom hardware architecture that efficiently executes AI workloads. The proposed accelerator focuses on two primary features:

- 1) Reconfigurable Processing Elements (PEs) for flexible operation mapping.
- 2) Caching and data reuse mechanisms to minimize external memory access.

This design aims to achieve high computational efficiency and energy savings, addressing the challenges in deploying AI models on resource-constrained platforms.

II. STATE-OF-THE-ART LITERATURE REVIEW

Several hardware platforms like GPUs, TPUs, FPGAs and ASICs are utilized to accelerate AI workloads, each optimized for specific performance, flexibility, and power trade-offs. These platforms are deployed in edge devices, data centers, and embedded systems for applications such as image recognition, natural language processing, and autonomous systems. Few recent hardware accelerators are discussed below:

A. Three-Dimensional GEMM Architecture with RISC-V [5]

An efficient General Matrix Multiplication (GEMM) module was proposed as the core of a Deep Learning Accelerator (DLA), integrating a RISC-V processor for scalar control and a processing engine for tensor operations as shown in Fig. 1. The GEMM module features 4096 MAC units supporting FP16 and INT8 computations with resource reuse through pipelined execution. The architecture employs a three-dimensional MAC arrangement for fine-grained control and high scalability via a hierarchical instruction queue. However, varying workload sizes may cause underutilization of resources, leading to the dark silicon effect.

B. Memory-Centric Modular Architecture for Neural Networks [6]

A modular Neural Processing Unit (NPU) architecture was introduced, consisting of 16 Processing Elements (PEs) and a




Fig. 1. Top Level architecture for DLA based on GEMM [5]




Fig. 2. Hardware configuration of Neural Engine for VGG-16 [6]

local SRAM to enable configurable parallelism and data reuse with architecture shown in Fig. 2. The design supports multiple convolution modes—depth-wise, point-wise, and 2D—through SRAM-based data configuration. By placing memory near the PEs, the design minimizes latency and power consumption. However, the fixed number of PEs may result in inefficiency for workloads that do not align with the hardware configuration.

C. Low-Power Architecture Exploiting Memory Hierarchy Optimization [7]

The row-stationary (RS) dataflow was proposed to reduce data movement energy cost in spatial architectures. This approach maximizes local data reuse of filter weights, activations, and partial sums across four storage hierarchy levels: DRAM, global buffer, inter-PE communication, and local register files. The design employs direct inter-PE communication and spatial parallelism to minimize high-cost DRAM accesses, significantly improving energy efficiency for CNN workloads.

D. Low-Power Architecture for Extreme Edge Inference [8]

TinyVers' FlexML accelerator targets ultra-low-power *tinyML* applications. It features an 8×8 SIMD array of precision-scalable MAC-based PEs, supporting runtime reconfigurable dataflow with zero-latency switching as shown in Fig. 3. The architecture enables symmetric precision quantization for activations and weights, minimizing hardware overhead




Fig. 3. Top Level architecture for DLA based on GEMM [8]

while maximizing data reuse. Output normalization through simple shift and ReLU operations further reduces computational cost. Although the design efficiently supports diverse ML workloads, mixed-precision handling may introduce bandwidth and control challenges.

III. PROPOSED METHODOLOGY

The hierarchy of proposed architecture begins with a “Processing Element (PE)”, and then comes the “Engine”, followed by a “Core” and then a “Mega Core” with it is a “Top Memory” and lastly the complete “Accelerator”. The bit precision(bp) of the accelerator is selected to be INT8 fixed for both weights and Input/Output feature values. The size of the input feature matrix is taken as 32×32 , while for weights it is 3×3 . As a proof of concept the architecture is designed to implement Le-Net5 Neural Network. The detailed explanation of sub blocks is given below:

A. Processing Element (PE)

The PE is basic workhorse in convolution operation and is responsible for the multiplication of feature and weight values. The inputs to PE are clock, enable, reset and weight value as shown in 4. However, the PE is unique in terms of its input feature. The PE has three inputs IF_{in1} , IF_{in2} , IF_{in3} for input feature instead of one. IF_{in1} is the input used immediately by the weight, IF_{in2} is used in 2nd iteration and IF_{in3} is used in third iteration of the weight feature multiplication. The value to be used in particular iteration depends on the value of “control” signal, which is also a 2-bit input to the PE. The PEs overall hardware consists of a Multiplier, a Mux and a register for storing the value at the particular iteration. The PE takes the the input with INT8 bit precision (bp) yielding the output value with INT16 bit precision, while taking three cycles to give output for one value of the input feature.

B. The Processing Engine

The next level in the hierarchy is the engine, with each engine consisting of two sets of 9 PEs as shown in 5. Each PE is connected to one of the 9 weights. The system assumes a weight stationary workload and the weights can only be changed when the engine is in “reset” or “not enable




Fig. 4. Processing Element




Fig. 5. Proposed Processing Engine

mode". Each set of 9 PEs has an adder tree which adds the outputs of the PEs together. Each engine is responsible for the convolution of a weight matrix with size $(kh,kw)=3 \times 3$ and input feature matrix with size $(Hs,Ws)=8 \times 8$. The engine has one feature buffer which is capable of storing 3 rows of (Hs,Ws) feature matrix. When the engine has its first three rows the operation starts. The engine is capable of producing the first row of 6×6 resulting matrix in 4 cycles. The engine also features an FSM-based control unit that generates the control input and dictates the operation of PEs and the engine. When the value of control is "0" the engine is allowed to take the value of next row of 8×8 input feature matrixes. On value "1" the PE units multiply the input IF_{in1} to the weight value. On value "2" the value at IF_{in2} is multiplied and on "3" the value at IF_{in3} is used. The connections to the three inputs for features are done via the input buffer with each input connected to its respective bit addressed value in buffer. The two sets of 9 PEs generate outputs in parallel.

C. The Processing Core

The level above the engine is the core. The core consists of one "Register file" with 64×8 bit dimensions. The register file is dual-port to allow pipelining between receiving data from the outside and feeding it to engine underneath. The Register file is chosen due to its lower area and better performance




Fig. 6. Proposed Processing Core

characteristics as compared to SRAM. The size of Register file is such that it is capable of holding an 8×8 feature matrix with a INT8 bit precision (bp). The addresses of the Reg file are generated by a small controller which is built with the register file. This controller provides a distributed control for each respective SRAM. This also produces addresses in fixed pattern since the engine processes the features in a fixed flow and not much flexibility is needed at this level for address generation. The component next to the Register file is the Engine (as described in previous subsection, which is fed the feature data from the Register file. The operation of the engine is enabled as soon as the Register file is fed the first three rows of feature 8×8 matrix. when the engine is ready for the next feature data it gives the feedback to the controller of register file, and the next value is fed to the engine from the register file. This allows pipelining between feeding data and processing it. The next unit in the core is the output "Buffer". The buffer is capable of holding 6×2 values of the output of engine i.e. two rows of the 6×6 output of the convolution results with INT16 precision. This buffer acts as input to the max-pooling, activation or batch normalization unit. Once the buffer is full, the maxpooling, activation or batch normalization unit is enabled to perform the relevant operation. Overall the core takes an 8×8 sub-section of the 32×32 input feature map along with the 3×3 weight vector. The result coming out of the core is a 3×3 max-pooled output in case of LeNet-5. This fusion of the convolution and pooling step allows for lesser hardware and lesser latency since applying pooling on a smaller matrix can be parallelized without the need for adding too much hardware. Since the pooling is parallelized on the 6×2 value buffer, the output is generated in one cycle and there is no need to iterate a large matrix as would be the case of traditional maxpooling on the complete 30×30 output of a traditional Le-Net 1st layer convolution. The overall architecture of the core is shown below in Fig. 6.

D. The Mega Core(Overall Accelerator)

The mega core takes a 256 bit vector as an input i.e. a 32×1 vector with an INT8 precision. This is provided by the SRAM on top of the mega core. This 256 bit vector is used to fill up the internal register files of the cores. The mega core also takes a 3×3 weight matrix with an INT8 precision as input. This is in fact a 72 bit vector with 9 values and




Fig. 7. Proposed Accelerator (The Mega Core)

each as 8 bits. The mega core also has a control mechanism implemented as a simple state machine counter combination to determine which core is filled up with the incoming input of 256 bit feature vector. The mega core has 5 buffers which store outputs of the results of the cores and give them at the output. Each buffer is 240 bits long i.e. can store 15 values with INT16 precision. This is effectively one row of the output 15x15 matrix with INT16 precision. The mega core consists of 25 cores each with its own register file, control unit for register file, convolution engine, buffer and maxpooling, activation and batch normalization unit. Each mega core is divided into five rows of five cores. Each row has a buffer which is capable of storing a 255 bit value. The segregation is done such that the effective stride of convolution matrix is one i.e., the last two values of the core1 are first two values of core2 in row 1 of the mega core. The cores within the mega core are all assigned the same weight 3x3 matrices. This ensures the weight stationary behavior and hence achieves parallelism in computing the results. The output of the mega core are the five buffers from each row of the core. The output is the vector of dimension 240(15x16)x1. Each row produces results at one cycle and no two rows produce results simultaneously. This is due to the fact that the data input to the mega core is one row at a time and the individual rows get filled in a round robin fashion and not continuously. This delay in output enables the result to be added to the results of the other kernels and stored to the output SRAM. Each buffer changes value three times for each 32x32x1 input. This results in a 15x15 matrix output. This output is stored in the output sram and can be transferred outside the chip.

IV. RESULTS AND DISCUSSIONS

The design was synthesized using Sky130nm technology which is a legacy node but the Process Design Kit (PDK) is open source. The whole process was carried out using open source toolchains Yosys, Openlane and OpenRAM. The clock frequency for the said design was 100MHz. The memory modules were placed as hard macros in the design and their characterization was done at Typical (TT) process corner. A brief comparison of the proposed accelerator design with state-of-the-art implementations is given in Table I.

The results show that the value of “GOPS per Watt” increases by almost 65% in the proposed design. This increase

TABLE I
COMPARISON WITH STATE-OF-THE-ART

Parameter	JSSC'23 [8]	This Work
Technology	Skywater 130 nm	Skywater 130 nm
Clock Frequency	100 MHz	100 MHz
Precision	INT8	INT8
GOPS (peak)	11.7	90
Power	46.8 mW	218 mW
GOPS/W	250	412.844

is due to the increases parallelization in the proposed design as compared to the JSSC'23 design. This increase in parallelization does increase power consumption but also increases the throughput by a much larger factor resulting in an overall increase in power efficiency.

V. CONCLUSION

The proposed work presents a hardware architecture optimized for accelerating diverse deep learning workloads by efficiently executing compute-intensive layers found across modern neural networks. The implemented accelerator demonstrates a 65% improvement over state-of-the-art design. These results highlight the suitability of the architecture for low-power, high-throughput edge-AI applications. While the current design offers competitive efficiency and performance, further enhancements can broaden its applicability and flexibility. Future improvements may include incorporating variable-precision computation within the existing tiling and dataflow structure to boost energy efficiency across different workloads.

REFERENCES

- [1] S. Das, S. Bhowmick, D. Burman, T. Roy, A. Bhar, and S. Bhattacharyya, “Optimizing cnn architectures for cat and dog classification with focus on generalization and robustness,” in *2025 International Conference on Computing, Intelligence, and Application (CIACON)*, 2025, pp. 1–6.
- [2] M. F. Tolba, H. Saleh, B. Mohammad, M. Al-Qutayri, and T. Stouraitis, “Enhanced cnn performance without retraining via weight approximation and data reuse,” in *2025 IEEE International Symposium on Circuits and Systems (ISCAS)*, 2025, pp. 1–5.
- [3] S. Swati, S. Kawa, R. Patel, A. M. Amrsh, M. S. Nagar, and P. Engineer, “Exploring quantization approaches for optimized training and inference for edge ai applications,” in *2025 11th International Conference on Communication and Signal Processing (ICCSP)*, 2025, pp. 1362–1367.
- [4] M. Usman, A. Zahid, and F. U. Din Farrukh, “Efficient multipliers for cnn with optimized compression techniques,” in *2023 20th International Bhurban Conference on Applied Sciences and Technology (IBCAST)*, 2023, pp. 291–296.
- [5] G. Lu, D. Xu, N. Wang, X. Zhang, D. Zhen, H. Lei, Y. Bai, D. Kong, H. Ruan, Z. Chi, X. Xiong, and K. Xu, “A design of 16tops efficient gemm module in deep learning accelerator,” in *2020 IEEE International Conference on Integrated Circuits, Technologies and Applications (ICTA)*, 2020, pp. 59–60.
- [6] E. Lee, M. Sung, S.-J. Jang, J. Park, and S.-S. Lee, “Memory-centric architecture of neural processing unit for edge device,” in *2021 18th International SoC Design Conference (ISOCC)*, 2021, pp. 240–241.
- [7] Y.-H. Chen, T.-J. Yang, J. Emer, and V. Sze, “Eyeriss v2: A flexible accelerator for emerging deep neural networks on mobile devices,” *IEEE Journal on Emerging and Selected Topics in Circuits and Systems*, vol. 9, no. 2, pp. 292–308, 2019.
- [8] V. Jain, S. Giraldo, J. D. Roose, L. Mei, B. Boons, and M. Verhelst, “Tinyvers: A tiny versatile system-on-chip with state-retentive emram for ml inference at the extreme edge,” *IEEE Journal of Solid-State Circuits*, vol. 58, no. 8, pp. 2360–2371, 2023.